

stepwise regression process to find best var

Steven

2024-02-23

Data source and original descriptions

data preparation and cleaning

Load necessary Libraries

```
library(tidyverse)
```

```
## — Attaching core tidyverse packages ————— tidyverse  
rse 2.0.0 —
```

```
## ✓ dplyr      1.1.1      ✓ readr      2.1.5
```

```
## ✓ forcats   1.0.0      ✓ stringr    1.5.0
```

```
## ✓ ggplot2   3.4.4      ✓ tibble     3.2.1
```

```
## ✓ lubridate 1.9.2      ✓ tidyr      1.3.0
```

```
## ✓ purrr     1.0.1
```

```
## — Conflicts ————— tidyverse_co  
nflicts() —
```

```
## ✘ dplyr::filter() masks stats::filter()
```

```
## ✘ dplyr::lag()     masks stats::lag()
```

```
## i Use the conflicted package (<http://conflicted.r-lib.org/>) to force all conflicts to become errors
```

```
if (!requireNamespace("caret", quietly = TRUE)) {  
  install.packages("caret")  
}
```

```
library(caret)
```

```
## 载入需要的程辑包: lattice
```

```
##
```

```
## 载入程辑包: 'caret'
```

```
##
```

```
## The following object is masked from 'package:purrr':
```

```
##
```

```
## lift
```

Load the data

```
data <- read_csv(file.choose()) # Open and Load file
```

```
## Warning: One or more parsing issues, call `problems()` on your data  
frame for details,
```

```
## e.g.:
```

```

## dat <- vroom(...)
## problems(dat)

## Rows: 34857 Columns: 21
## — Column specification —————
## Delimiter: ","
## chr (8): Suburb, Address, Type, Method, SellerG, Date, CouncilArea,
Regionname
## dbl (13): Rooms, Price, Distance, Postcode, Bedroom2, Bathroom, Car,
Landsiz...
##
## i Use `spec()` to retrieve the full column specification for this da
ta.
## i Specify the column types or set `show_col_types = FALSE` to quiet
this message.

# Correct column names
names(data)[names(data) == "Lattitude"] <- "Latitude"
names(data)[names(data) == "Longtitude"] <- "Longitude"

# Remove unnecessary columns using dplyr's select function
data_clean <- data %>%
  dplyr::select(Suburb, Rooms, Type, Price, Distance, Bedroom2, Bathroo
m, Car, Landsize, BuildingArea, YearBuilt, CouncilArea, Latitude, Longi
tude, Propertycount, Date)

# Convert 'Date' to date type
data_clean$Date <- as.Date(data_clean$Date, format = "%d/%m/%Y")

# Calculate 'YearsAfterBuilt'
data_clean$YearsAfterBuilt <- as.numeric(format(data_clean$Date, "%Y"))
- data_clean$YearBuilt

# Calculate PricePerBuildingArea
data_clean$PricePerBuildingArea <- data_clean$Price / data_clean$Buildi
ngArea

# Remove rows with missing values
data_clean <- na.omit(data_clean)

# Drop the "Price", "Longitude", "Latitude", "YearBuilt" and columns fr
om the dataset
data_clean <- subset(data_clean, select = -c(Price, Longitude, Latitude
, YearBuilt))

# Convert categorical variables to factors
cat_vars <- c("Suburb", "Type", "CouncilArea") # Add categorical varia

```

bles here

```
data_clean[cat_vars] <- lapply(data_clean[cat_vars], as.factor)

# Convert non-categorical variables to numeric
non_cat_vars <- setdiff(names(data_clean), c(cat_vars, "PricePerBuildingArea"))
data_clean[non_cat_vars] <- lapply(data_clean[non_cat_vars], as.numeric)

# Standardize non-categorical variables
data_clean[non_cat_vars] <- scale(data_clean[non_cat_vars])

# Separate predictors and target variable
predictors <- setdiff(names(data_clean), "PricePerBuildingArea")

# Split data into training and testing sets
set.seed(123)
indexes <- createDataPartition(data_clean$PricePerBuildingArea, p = 0.8,
, list = FALSE)
train_data <- data_clean[indexes, ]
test_data <- data_clean[-indexes, ]
```

Model training and AIC process

```
# Remove rows with NA, NaN, or Inf values in the target variable
train_data <- train_data[!is.na(train_data$PricePerBuildingArea) & !is.nan(
train_data$PricePerBuildingArea) & !is.infinite(train_data$PricePerBuildingArea), ]

# Train stepwise regression model
model <- step(lm(PricePerBuildingArea ~ ., data = train_data[, c(predictors, "PricePerBuildingArea")]), direction = "backward")

## Start: AIC=146189.8
## PricePerBuildingArea ~ Suburb + Rooms + Type + Distance + Bedroom2 +
## Bathroom + Car + Landsize + BuildingArea + CouncilArea +
## Propertycount + Date + YearsAfterBuilt
##
##
## Step: AIC=146189.8
## PricePerBuildingArea ~ Suburb + Rooms + Type + Distance + Bedroom2 +
## Bathroom + Car + Landsize + BuildingArea + CouncilArea +
## Date + YearsAfterBuilt
##
##
## Step: AIC=146189.8
## PricePerBuildingArea ~ Suburb + Rooms + Type + Distance + Bedroom2 +
```

```

## Bathroom + Car + Landsize + BuildingArea + Date + YearsAfterBuilt
##
##           Df Sum of Sq      RSS      AIC
## - Suburb   305 2.1970e+11 6.3092e+12 145831
## - Landsize    1 4.9944e+07 6.0895e+12 146188
## - Bathroom    1 7.9612e+07 6.0896e+12 146188
## - Date        1 5.0891e+08 6.0900e+12 146188
## - YearsAfterBuilt 1 1.0302e+09 6.0905e+12 146189
## <none>                                6.0895e+12 146190
## - Car        1 2.9181e+09 6.0924e+12 146191
## - Rooms      1 3.3703e+09 6.0928e+12 146192
## - Distance   1 4.5451e+09 6.0940e+12 146193
## - Bedroom2   1 1.0378e+10 6.0999e+12 146200
## - Type       2 4.0714e+10 6.1302e+12 146233
## - BuildingArea 1 1.1391e+11 6.2034e+12 146319
##
## Step: AIC=145830.6
## PricePerBuildingArea ~ Rooms + Type + Distance + Bedroom2 + Bathroom
+
## Car + Landsize + BuildingArea + Date + YearsAfterBuilt
##
##           Df Sum of Sq      RSS      AIC
## - Date        1 1.5248e+07 6.3092e+12 145829
## - Landsize    1 7.3854e+07 6.3092e+12 145829
## <none>                                6.3092e+12 145831
## - Bathroom    1 2.5749e+09 6.3117e+12 145831
## - Car         1 3.0252e+09 6.3122e+12 145832
## - Rooms       1 4.0337e+09 6.3132e+12 145833
## - Bedroom2    1 9.7407e+09 6.3189e+12 145839
## - YearsAfterBuilt 1 1.0803e+10 6.3200e+12 145841
## - Distance    1 1.3718e+10 6.3229e+12 145844
## - Type        2 4.0314e+10 6.3495e+12 145872
## - BuildingArea 1 9.5400e+10 6.4046e+12 145935
##
## Step: AIC=145828.6
## PricePerBuildingArea ~ Rooms + Type + Distance + Bedroom2 + Bathroom
+
## Car + Landsize + BuildingArea + YearsAfterBuilt
##
##           Df Sum of Sq      RSS      AIC
## - Landsize    1 7.6641e+07 6.3093e+12 145827
## <none>                                6.3092e+12 145829
## - Bathroom    1 2.5787e+09 6.3118e+12 145829
## - Car         1 3.0414e+09 6.3122e+12 145830
## - Rooms       1 4.1124e+09 6.3133e+12 145831
## - Bedroom2    1 9.9601e+09 6.3191e+12 145838
## - YearsAfterBuilt 1 1.0788e+10 6.3200e+12 145839
## - Distance    1 1.4390e+10 6.3236e+12 145843

```

```

## - Type          2 4.0321e+10 6.3495e+12 145870
## - BuildingArea  1 9.5635e+10 6.4048e+12 145933
##
## Step: AIC=145826.7
## PricePerBuildingArea ~ Rooms + Type + Distance + Bedroom2 + Bathroom
+
##   Car + BuildingArea + YearsAfterBuilt
##
##           Df  Sum of Sq      RSS    AIC
## <none>
##           1 2.5938e+09 6.3119e+12 145827
## - Bathroom
##           1 3.1298e+09 6.3124e+12 145828
## - Car
##           1 4.1043e+09 6.3134e+12 145829
## - Rooms
##           1 9.9565e+09 6.3192e+12 145836
## - Bedroom2
##           1 1.0813e+10 6.3201e+12 145837
## - YearsAfterBuilt
##           1 1.4324e+10 6.3236e+12 145841
## - Distance
##           2 4.0248e+10 6.3495e+12 145868
## - Type
##           1 9.5562e+10 6.4048e+12 145931

# Make predictions on test data
predictions <- predict(model, newdata = test_data)

# Evaluate the model
rmse <- sqrt(mean((predictions - test_data$PricePerBuildingArea)^2))
print(paste("RMSE: ", rmse))

## [1] "RMSE:  Inf"

```

Displaying AIC value graph

```

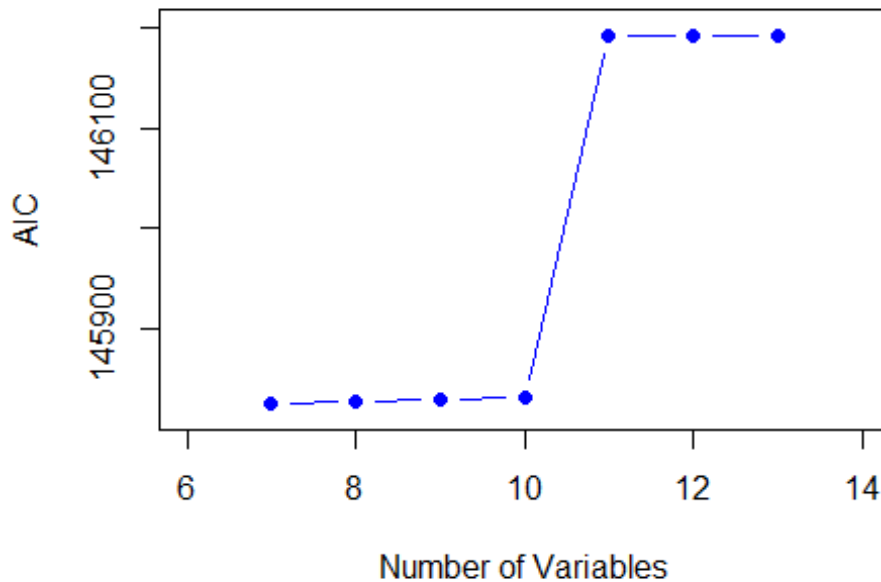
# AIC values from the stepwise regression process
aic_values <- c(146191.6, 146191.6, 146191.6, 145831.9, 145830, 145828.6, 145826.6)

# Number of variables in each step
num_variables <- c(13, 12, 11, 10, 9, 8, 7)

# Plotting the AIC values against the number of variables
plot(num_variables, aic_values, type = "b",
     xlab = "Number of Variables",
     ylab = "AIC",
     main = "Stepwise Regression: AIC vs. Number of Variables",
     xlim = c(min(num_variables) - 1, max(num_variables) + 1),
     ylim = c(min(aic_values) - 10, max(aic_values) + 10),
     col = "blue",
     pch = 19)

```

Stepwise Regression: AIC vs. Number of Variable



final model summary

```
# Train the final model based on the selected predictors
final_model <- lm(PricePerBuildingArea ~ Rooms + Type + Distance + Bedroom2 +
                 Bathroom + Car + BuildingArea + YearsAfterBuilt, data = train_data)
```

```
# Print the summary of the final model
summary(final_model)
```

```
##
## Call:
## lm(formula = PricePerBuildingArea ~ Rooms + Type + Distance +
##     Bedroom2 + Bathroom + Car + BuildingArea + YearsAfterBuilt,
##     data = train_data)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -31541  -3928  -1932    471 1034369
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    8685.3     451.6  19.233 < 2e-16 ***
## Rooms         -3085.5     1439.3  -2.144  0.032083 *
## Typet          9392.8     1455.1   6.455  1.15e-10 ***
```

```

## Typeu          461.1      1229.1    0.375 0.707573
## Distance       -1655.3      413.3   -4.005 6.27e-05 ***
## Bedroom2       4794.0     1435.7    3.339 0.000845 ***
## Bathroom        851.7      499.7    1.704 0.088379 .
## Car             745.5      398.2    1.872 0.061235 .
## BuildingArea   -4803.9      464.4  -10.344 < 2e-16 ***
## YearsAfterBuilt 1495.5      429.8    3.480 0.000505 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 29880 on 7065 degrees of freedom
## Multiple R-squared:  0.02958,    Adjusted R-squared:  0.02834
## F-statistic: 23.93 on 9 and 7065 DF,  p-value: < 2.2e-16

```

variable correlation check

```

if (!requireNamespace("GGally", quietly = TRUE)) {
  install.packages("GGally")
}

## Registered S3 method overwritten by 'GGally':
##   method from
##   +.gg   ggplot2

# Load the GGally package
library(GGally)

# Select predictors for correlation analysis
cor_data <- train_data[, c("Rooms", "Type", "Distance", "Bedroom2", "Bathroom", "Car", "BuildingArea", "YearsAfterBuilt")]

# Convert non-numeric columns to numeric
cor_data_numeric <- as.data.frame(sapply(cor_data, as.numeric))

# Compute pairwise correlations
correlation_matrix <- cor(cor_data_numeric)

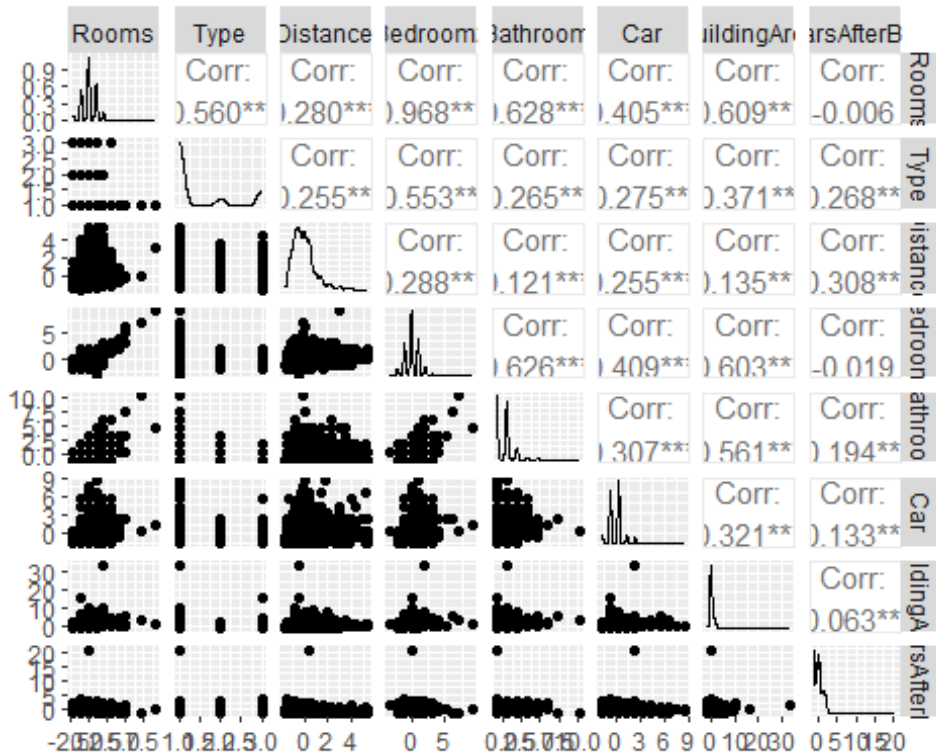
# Print pairwise correlations
print(correlation_matrix)

##           Rooms      Type  Distance  Bedroom2  Bathroom
## Rooms      1.00000000 -0.5597444  0.2800834  0.96801199  0.6276422
## Type       -0.559744386  1.0000000 -0.2545656 -0.55332318 -0.2652077
## Distance   0.280083407 -0.2545656  1.0000000  0.28843088  0.1209617
## Bedroom2   0.968011994 -0.5533232  0.2884309  1.00000000  0.6261735

```

```
## Bathroom      0.627642191 -0.2652077  0.1209617  0.62617354  1.00
00000
## Car            0.405160513 -0.2754376  0.2553177  0.40921543  0.30
73153
## BuildingArea  0.609242417 -0.3708441  0.1348149  0.60297412  0.56
13127
## YearsAfterBuilt -0.006047909 -0.2679670 -0.3078093 -0.01937879 -0.19
36693
##              Car BuildingArea YearsAfterBuilt
## Rooms        0.4051605  0.60924242  -0.006047909
## Type         -0.2754376 -0.37084408  -0.267966986
## Distance     0.2553177  0.13481485  -0.307809292
## Bedroom2    0.4092154  0.60297412  -0.019378790
## Bathroom     0.3073153  0.56131273  -0.193669268
## Car          1.0000000  0.32071519  -0.133360609
## BuildingArea 0.3207152  1.00000000  -0.063251087
## YearsAfterBuilt -0.1333606 -0.06325109  1.000000000
```

```
# Create a histogram grid for visualization
ggpairs(cor_data_numeric)
```



correlation graph display

```
# Load necessary Libraries
```

```
library(corrplot)
```

```
## corrplot 0.92 loaded
```



```

# Convert non-numeric columns to numeric
cor_data_numeric <- as.data.frame(sapply(cor_data, as.numeric))

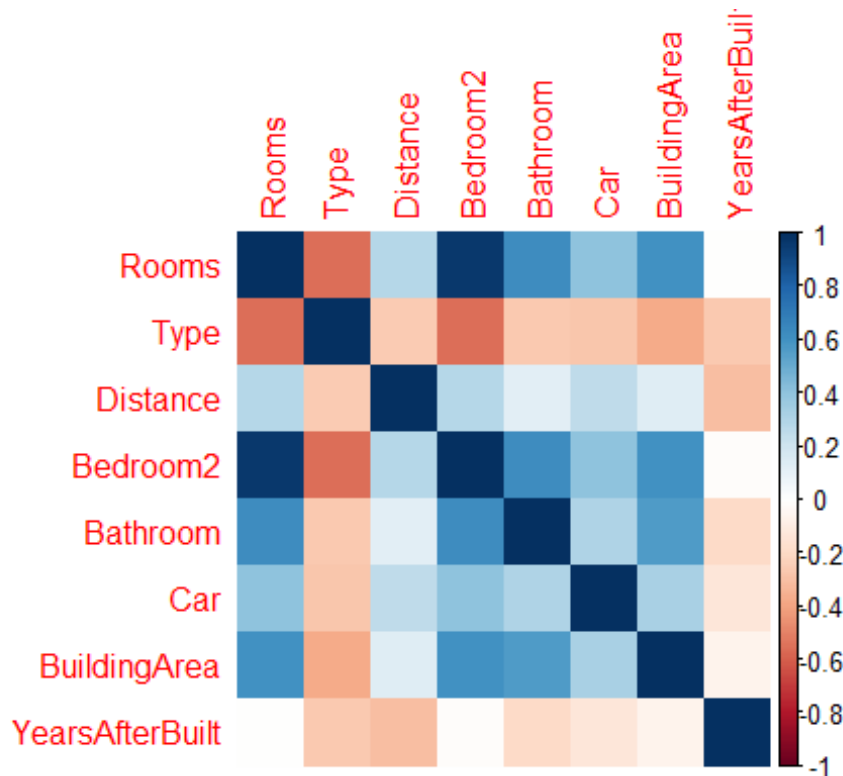
# Compute pairwise correlations
correlation_matrix <- cor(cor_data_numeric)

# Print pairwise correlations
print(correlation_matrix)

##              Rooms      Type  Distance  Bedroom2  Bathroom
## Rooms      1.00000000 -0.5597444  0.2800834  0.96801199  0.6276422
## Type      -0.559744386  1.0000000 -0.2545656 -0.55332318 -0.2652077
## Distance   0.280083407 -0.2545656  1.0000000  0.28843088  0.1209617
## Bedroom2   0.968011994 -0.5533232  0.2884309  1.00000000  0.62617354
## Bathroom   0.627642191 -0.2652077  0.1209617  0.62617354  1.0000000
## Car        0.405160513 -0.2754376  0.2553177  0.40921543  0.3073153
## BuildingArea 0.609242417 -0.3708441  0.1348149  0.60297412  0.56131273
## YearsAfterBuilt -0.006047909 -0.2679670 -0.3078093 -0.01937879 -0.193669268
##              Car BuildingArea YearsAfterBuilt
## Rooms      0.4051605  0.60924242  -0.006047909
## Type      -0.2754376 -0.37084408  -0.267966986
## Distance   0.2553177  0.13481485  -0.307809292
## Bedroom2   0.4092154  0.60297412  -0.019378790
## Bathroom   0.3073153  0.56131273  -0.193669268
## Car        1.0000000  0.32071519  -0.133360609
## BuildingArea 0.3207152  1.00000000  -0.063251087
## YearsAfterBuilt -0.1333606 -0.06325109  1.000000000

# Create a correlation plot with color
corrplot(correlation_matrix, method = "color")

```



Actual vs predicted graph

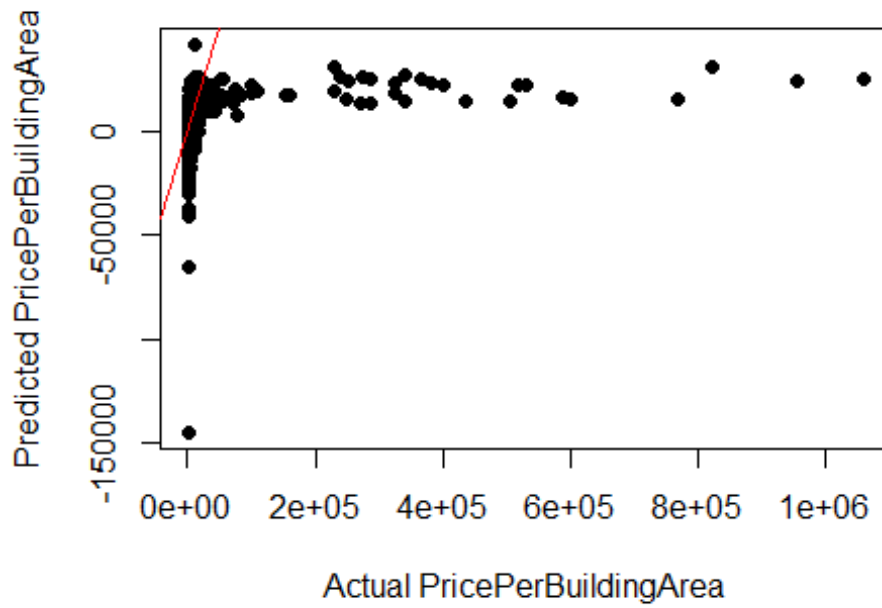
```
# Calculate predicted values using the model
predicted_values <- predict(model, newdata = train_data)

# Ensure that only rows with no missing values in the relevant columns
# are used
valid_rows <- complete.cases(train_data[c("PricePerBuildingArea", "Rooms",
"Type", "Distance", "Bedroom2", "Bathroom", "Car", "BuildingArea",
"YearsAfterBuilt")])
actual_values <- train_data$PricePerBuildingArea[valid_rows]
predicted_values <- predicted_values[valid_rows]

# Now plot actual vs predicted values
plot(actual_values, predicted_values,
      main = "Actual vs Predicted PricePerBuildingArea",
      xlab = "Actual PricePerBuildingArea",
      ylab = "Predicted PricePerBuildingArea",
      pch = 19) # pch = 19 for solid circles

# Add a Line of perfect fit
abline(a = 0, b = 1, col = "red")
```

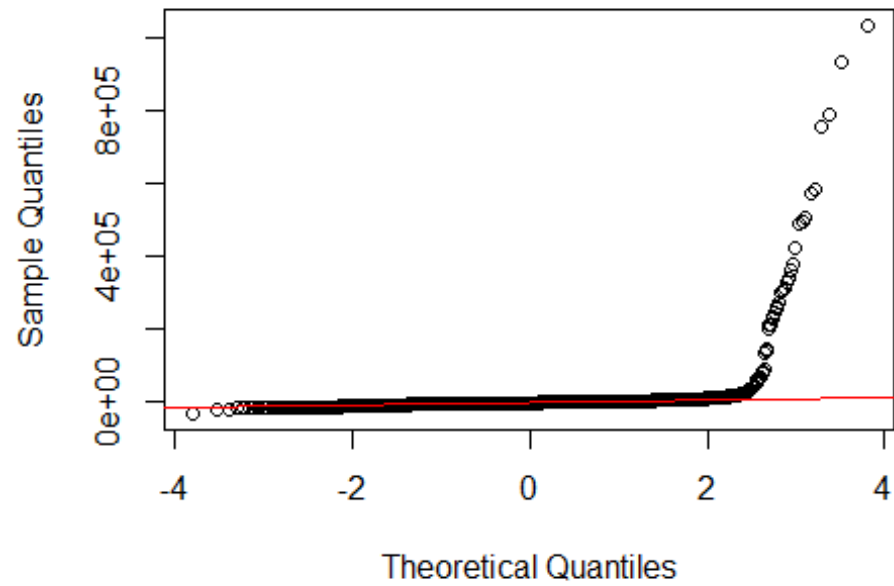
Actual vs Predicted PricePerBuildingArea



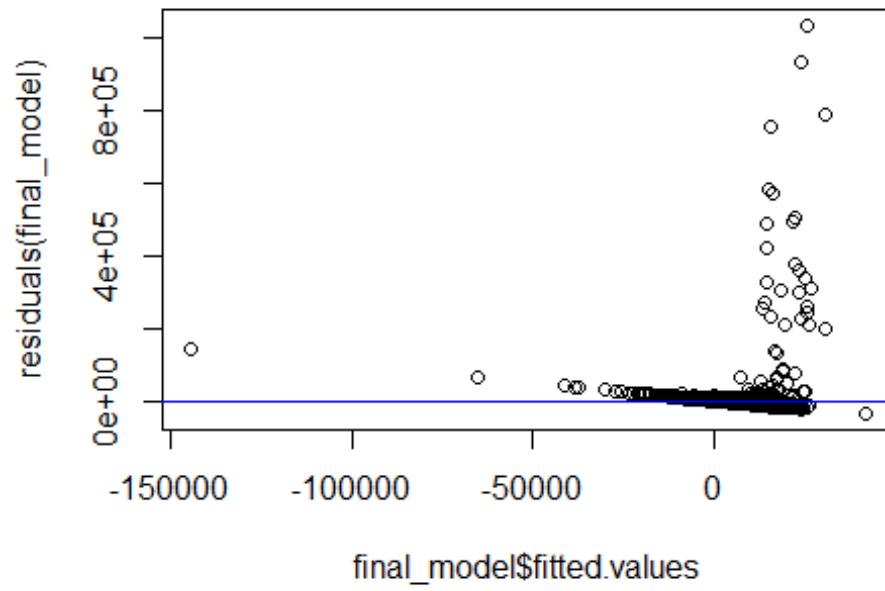
qq and residual plots

```
# QQ plot for the first model  
qqnorm(residuals(final_model))  
qqline(residuals(final_model), col = "red")
```

Normal Q-Q Plot



```
# Residual plot for the first model  
plot(final_model$fitted.values, residuals(final_model))  
abline(h = 0, col = "blue")
```



of results # Comment of business implications

Comment