

# classification, regression with outlier exclusion

Steven

2024-02-16

## Data source and original descriptions

### Step 1: Data Cleaning and Preparation

```
# Load necessary Libraries
```

```
library(dplyr)
```

```
##
```

```
## 载入程辑包: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
##   filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
##   intersect, setdiff, setequal, union
```

```
library(readr)
```

```
# Step 1: Load the dataset
```

```
data_1 <- read_csv(file.choose()) # select the file interactively
```

```
## Warning: One or more parsing issues, call `problems()` on your data  
## frame for details,
```

```
## e.g.:
```

```
##   dat <- vroom(...)
```

```
##   problems(dat)
```

```
## Rows: 34857 Columns: 21
```

```
## — Column specification
```

---

```
## Delimiter: ","
```

```
## chr (8): Suburb, Address, Type, Method, SellerG, Date, CouncilArea,  
## Regionname
```

```
## dbl (13): Rooms, Price, Distance, Postcode, Bedroom2, Bathroom, Car,  
## Landsiz...
```

```
##
```

```
## i Use `spec()` to retrieve the full column specification for this  
## data.
```

```

## i Specify the column types or set `show_col_types = FALSE` to quiet
this message.

# Step 2: Correct typos in column names
names(data_1)[names(data_1) == "Lattitude"] <- "Latitude"
names(data_1)[names(data_1) == "Longtitude"] <- "Longitude"

# Step 3: Calculate features
# Convert 'Date' from character to Date type.

data_1$Date <- as.Date(data_1$Date, format = "%d/%m/%Y")

# Extract the year from the 'Date' column
data_1$YearOfSale <- as.numeric(format(data_1$Date, "%Y"))

# Calculate 'YearsSinceBuilt' and 'Priceperbuildingarea'
data_1$YearsSinceBuilt <- data_1$YearOfSale - data_1$YearBuilt
data_1$Priceperbuildingarea <- with(data_1, Price / BuildingArea)

# Step 4: Clean the data (Eliminate `NA` and `Inf` values)
data_1 <- na.omit(data_1) # Remove rows with NA values

#Identify numeric columns
numeric_cols <- sapply(data_1, is.numeric)

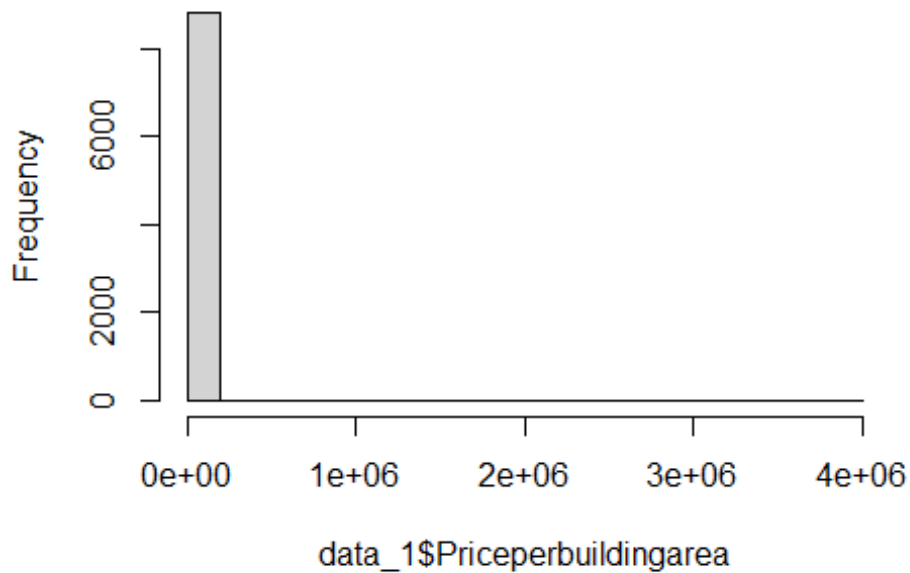
# Apply is.infinite only to numeric columns and then reduce to rows
with any Inf values
rows_with_inf <- apply(data_1[, numeric_cols], 1, function(x)
any(is.infinite(x)))

# Remove rows with Inf values
data_1 <- data_1[!rows_with_inf, ]

# hist 1
hist(data_1$Priceperbuildingarea)

```

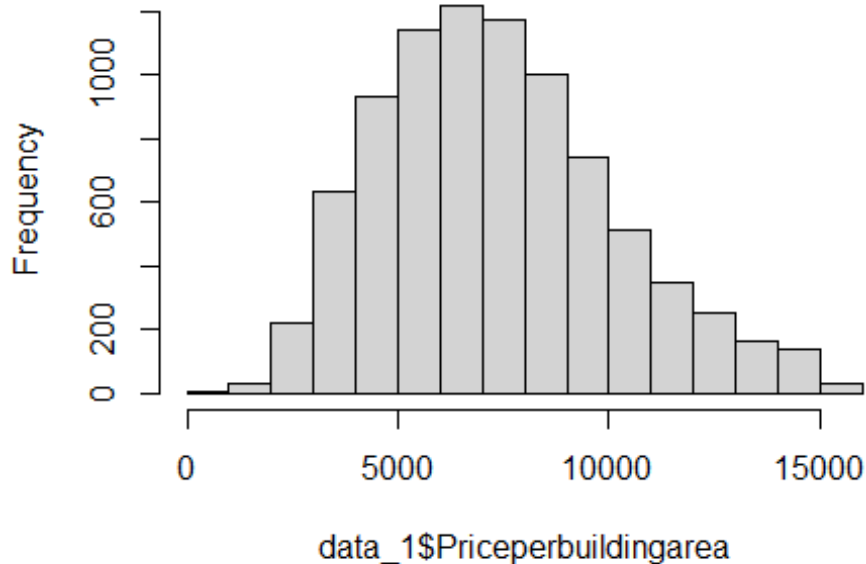
## Histogram of data\_1\$Priceperbuildingarea



```
# Step 5: Eliminate outliers in 'Priceperbuildingarea'
Q1 <- quantile(data_1$Priceperbuildingarea, 0.25, na.rm = TRUE)
Q3 <- quantile(data_1$Priceperbuildingarea, 0.75, na.rm = TRUE)
IQR <- Q3 - Q1
data_1 <- data_1 %>%
  filter(Priceperbuildingarea >= (Q1 - 1.5 * IQR) &
Priceperbuildingarea <= (Q3 + 1.5 * IQR))

# hist 2
hist(data_1$Priceperbuildingarea)
```

## Histogram of data\_1\$Priceperbuildingarea



```
# Using names() function
column_names <- names(data_1)
print(column_names)

## [1] "Suburb"           "Address"          "Rooms"
## [4] "Type"            "Price"            "Method"
## [7] "SellerG"         "Date"             "Distance"
## [10] "Postcode"        "Bedroom2"         "Bathroom"
## [13] "Car"             "Landsize"         "BuildingArea"
## [16] "YearBuilt"       "CouncilArea"      "Latitude"
## [19] "Longitude"       "Regionname"       "Propertycount"
## [22] "YearOfSale"     "YearsSinceBuilt"
"Priceperbuildingarea"

# Identify numeric columns
numeric_columns <- sapply(data_1, is.numeric)

# Select only numeric columns
selected_data <- data_1[, numeric_columns]

# Calculate correlation matrix
correlation_matrix <- cor(selected_data, use = "pairwise.complete.obs")

# Print correlation matrix
print(correlation_matrix)
```

##	Rooms	Price	Distance
Postcode			
## Rooms	1.0000000000	0.501094726	0.279584628
0.08667307			
## Price	0.5010947257	1.000000000	-0.220490168
0.03541692			
## Distance	0.2795846276	-0.220490168	1.000000000
0.50293962			
## Postcode	0.0866730687	0.035416925	0.502939623
1.00000000			
## Bedroom2	0.9642190596	0.485052942	0.286713808
0.08980185			
## Bathroom	0.6218748893	0.488440346	0.124599926
0.11424239			
## Car	0.4021756258	0.216086293	0.260624274
0.06013886			
## Landsize	0.0990786454	0.058431506	0.138053610
0.07187845			
## BuildingArea	0.6245141720	0.561318543	0.136755737
0.08186941			
## YearBuilt	0.0007745297	-0.295673010	0.297603488
0.10279907			
## Latitude	0.0175199666	-0.224470815	-0.066282456 -
0.19444725			
## Longitude	0.0822794916	0.224735399	0.165473824
0.36046687			
## Propertycount	-0.0836318090	-0.063615943	-0.003382938
0.03570636			
## YearOfSale	0.1913029831	-0.002704971	0.327247947
0.12848419			
## YearsSinceBuilt	0.0026839596	0.296210171	-0.292274590 -
0.10067900			
## Priceperbuildingarea	-0.1837209402	0.513547389	-0.514936406 -
0.05173650			
##	Bedroom2	Bathroom	Car
Landsize			
## Rooms	0.964219060	0.62187489	0.40217563
0.099078645			
## Price	0.485052942	0.48844035	0.21608629
0.058431506			
## Distance	0.286713808	0.12459993	0.26062427
0.138053610			
## Postcode	0.089801854	0.11424239	0.06013886
0.071878455			
## Bedroom2	1.000000000	0.62524398	0.40534655
0.099037744			
## Bathroom	0.625243979	1.00000000	0.31099793
0.075464539			
## Car	0.405346551	0.31099793	1.00000000
0.120537396			

## Landsize 1.000000000	0.099037744	0.07546454	0.12053740
## BuildingArea 0.081978978	0.613469920	0.57272213	0.32271057
## YearBuilt 0.034314694	0.011124211	0.19707483	0.13496208
## Latitude 0.042805418	0.021817673	-0.04430443	0.01646177
## Longitude 0.009794645	0.081787849	0.11170203	0.03390007 -
## Propertycount 0.032234320	-0.081790074	-0.06013173	-0.03065207 -
## YearOfSale 0.085526603	0.213542728	0.11221420	0.15686142
## YearsSinceBuilt 0.032835826	-0.007283995	-0.19543590	-0.13239251 -
## Priceperbuildingarea 0.040829514	-0.187782815	-0.13925201	-0.16947349 -
## Longitude	BuildingArea	YearBuilt	Latitude
## Rooms 0.082279492	0.62451417	0.0007745297	0.01751997
## Price 0.224735399	0.56131854	-0.2956730101	-0.22447081
## Distance 0.165473824	0.13675574	0.2976034883	-0.06628246
## Postcode 0.360466868	0.08186941	0.1027990711	-0.19444725
## Bedroom2 0.081787849	0.61346992	0.0111242112	0.02181767
## Bathroom 0.111702026	0.57272213	0.1970748307	-0.04430443
## Car 0.033900071	0.32271057	0.1349620821	0.01646177
## Landsize 0.009794645	0.08197898	0.0343146942	0.04280542 -
## BuildingArea 0.101937067	1.00000000	0.0623103272	-0.03368642
## YearBuilt 0.024585573	0.06231033	1.0000000000	0.08946128 -
## Latitude 0.347901002	-0.03368642	0.0894612827	1.00000000 -
## Longitude 1.00000000	0.10193707	-0.0245855731	-0.34790100
## Propertycount 0.025165982	-0.06024646	0.0170254136	0.02877662
## YearOfSale 0.020664425	0.08701842	0.1185121940	0.04888162
## YearsSinceBuilt 0.025008056	-0.06085997	-0.9998387213	-0.08875451

```

## Priceperbuildingarea -0.23622285 -0.5193228334 -0.26808344
0.186735548
##
## Propertycount YearOfSale YearsSinceBuilt
## Rooms -0.083631809 0.191302983 0.002683960
## Price -0.063615943 -0.002704971 0.296210171
## Distance -0.003382938 0.327247947 -0.292274590
## Postcode 0.035706359 0.128484185 -0.100678995
## Bedroom2 -0.081790074 0.213542728 -0.007283995
## Bathroom -0.060131728 0.112214200 -0.195435898
## Car -0.030652066 0.156861415 -0.132392512
## Landsize -0.032234320 0.085526603 -0.032835826
## BuildingArea -0.060246459 0.087018419 -0.060859970
## YearBuilt 0.017025414 0.118512194 -0.999838721
## Latitude 0.028776623 0.048881622 -0.088754510
## Longitude 0.025165982 0.020664425 0.025008056
## Propertycount 1.000000000 0.019590133 -0.016704842
## YearOfSale 0.019590133 1.000000000 -0.100660488
## YearsSinceBuilt -0.016704842 -0.100660488 1.000000000
## Priceperbuildingarea -0.006301737 -0.132182942 0.517961496
##
## Priceperbuildingarea
## Rooms -0.183720940
## Price 0.513547389
## Distance -0.514936406
## Postcode -0.051736503
## Bedroom2 -0.187782815
## Bathroom -0.139252010
## Car -0.169473491
## Landsize -0.040829514
## BuildingArea -0.236222853
## YearBuilt -0.519322833
## Latitude -0.268083437
## Longitude 0.186735548
## Propertycount -0.006301737
## YearOfSale -0.132182942
## YearsSinceBuilt 0.517961496
## Priceperbuildingarea 1.000000000


# Load the corrplot package
library(corrplot)

## corrplot 0.92 loaded

# Calculate the correlation matrix
correlation_matrix <- cor(selected_data, use = "pairwise.complete.obs")

# Create the colored correlation grid
corrplot(correlation_matrix, method = "color")

```




## Step 2: Classification with K-Means Clustering:

```
library(dplyr)
library(ggplot2)
library(cluster) # For clustering

# coordinates
coords <- data_1 %>% select(Longitude, Latitude)

# Determine the optimal number of clusters (optional, for illustration)
# This step can be computationally intensive for large datasets
wss <- (nrow(coords)-1)*sum(apply(coords,2,var))
for (i in 2:15) wss[i] <- sum(kmeans(coords, centers=i)$withinss)
plot(1:15, wss, type="b", xlab="Number of Clusters", ylab="Within
groups sum of squares")
```





```

# K-Means Clustering
set.seed(123) # For reproducibility
k <- 4 # Choose based on analysis, e.g., using the Elbow Method above
km_res <- kmeans(coords, centers = k)

# Assign class numbers to the original data and factorize
data_1$Class <- km_res$cluster
data_1$Class <- factor(data_1$Class)


# Step 4: Visualize on a Map
library(ggmap)

## i Google's Terms of Service: <https://mapsplatform.google.com>
## i Please cite ggmap if you use it! Use `citation("ggmap")` for
details.

library(ggplot2)

# Basic plot with ggplot2
ggplot(data_1, aes(x = Longitude, y = Latitude, color = factor(Class)))
+
  geom_point(alpha = 0.5) +
  labs(title = "Spatial Clustering of Data Points", color = "Class") +
  theme_minimal()

```



### Step 3: visualize histograms:

```
library(ggplot2)
library(dplyr)


# 'data_1' contains 'Y_label' and 'Class' columns
# Loop through each class and plot a histogram
unique_classes <- unique(data_1$Class)

# Create a list to store plots
plot_list <- list()

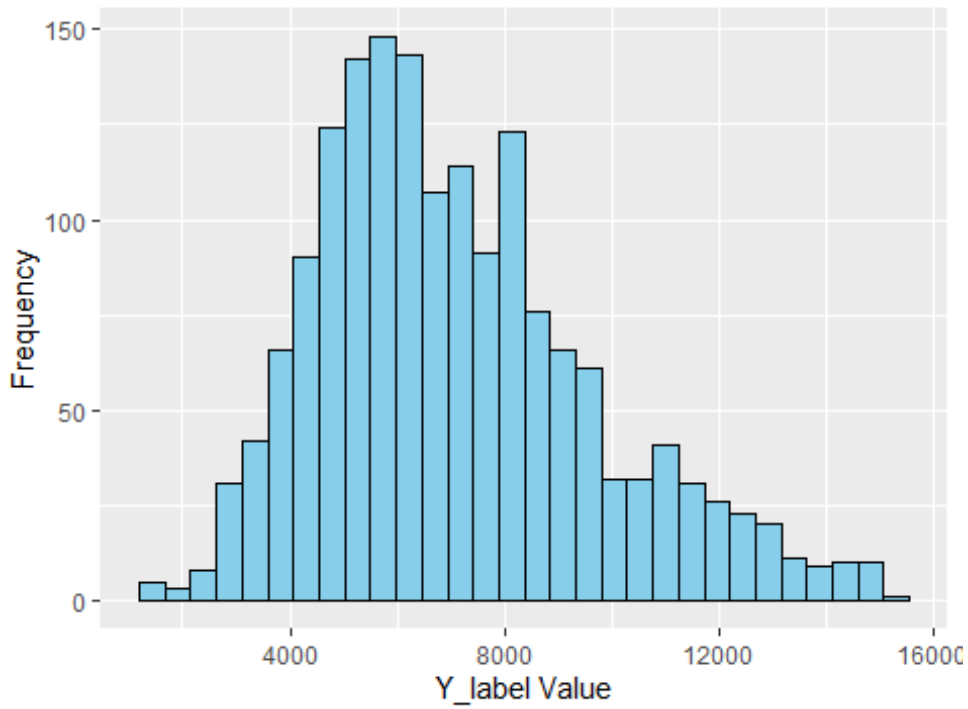
for(class in unique_classes) {
  plot <- data_1 %>%
    filter(Class == class) %>%
    ggplot(aes(x = Priceperbuildingarea)) +
    geom_histogram(bins = 30, fill = "skyblue", color = "black") +
    ggtitle(paste("Histogram of Priceperbuildingarea for Class",
class)) +
    xlab("Y_label Value") +
    ylab("Frequency")

  print(plot) # Display the plot
  plot_list[[as.character(class)]] <- plot # Store the plot in a list
}
```


### Histogram of Priceperbuildingarea for Class 2




### Histogram of Priceperbuildingarea for Class 1



Histogram of Priceperbuildingarea for Class 4



Histogram of Priceperbuildingarea for Class 3



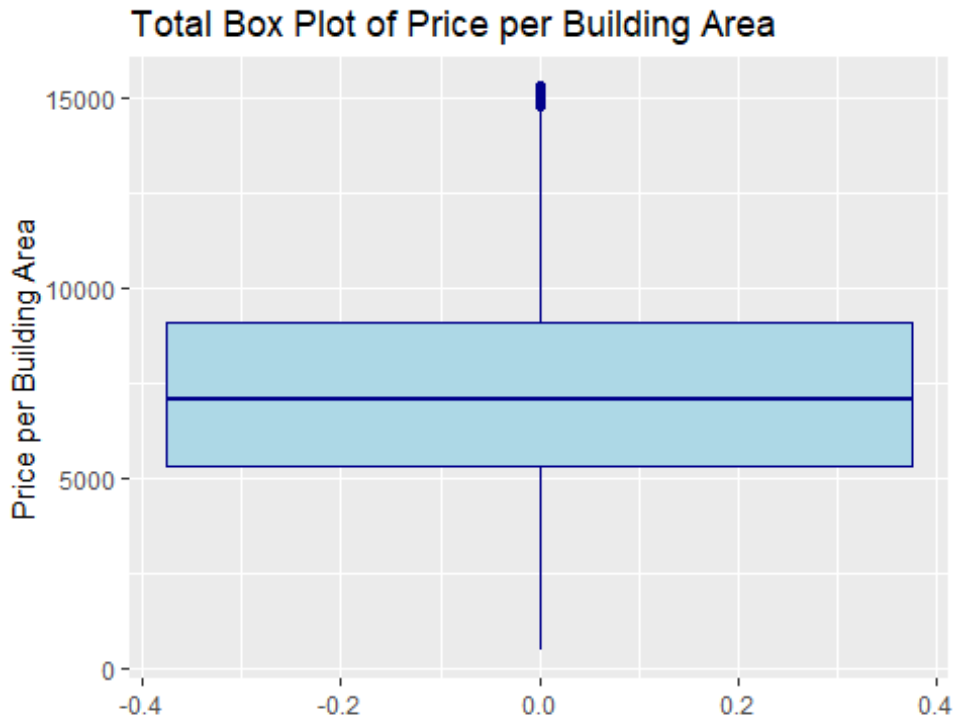
#### Step 4: Box plots

```
library(ggplot2)
```

```

# Total Box Plot for 'Priceperbuildingarea'
ggplot(data_1, aes(y = Priceperbuildingarea)) +
  geom_boxplot(fill = "lightblue", color = "darkblue") +
  ggtitle("Total Box Plot of Price per Building Area") +
  ylab("Price per Building Area") +
  xlab("")

```



```

# Box Plots for 'Priceperbuildingarea' by Class
ggplot(data_1, aes(x = factor(Class), y = Priceperbuildingarea, fill =
factor(Class))) +
  geom_boxplot() +
  scale_fill_brewer(palette = "Pastel1") + # Color scheme
  ggtitle("Box Plot of Price per Building Area by Class") +
  xlab("Class") +
  ylab("Price per Building Area") +
  theme_light() +
  theme(legend.title = element_blank()) # Remove the Legend title


```



#### Step 5: Visualization of Price distribution

```
library(ggplot2)  
library(ggmap)
```

```
# PricePerBuildingArea, visualize classification based on it  
ggplot(data_1, aes(x = Longitude, y = Latitude, color =  
Priceperbuildingarea)) + geom_point() + theme_minimal() +  
ggtitle("Price per Building Area Distribution")
```



## Step 6: Split Data into Training and Testing

```
library(caret)
```

```
## 载入需要的程辑包: lattice
```

```
set.seed(123) # For reproducibility
index <- createDataPartition(data_1$Priceperbuildingarea, p = 0.8, list
= FALSE)
trainData_1 <- data_1[index, ]
testData_1 <- data_1[-index, ]

trainData_trimmed_1 <- subset(trainData_1, select = c(Class,
YearsSinceBuilt, Priceperbuildingarea))
#trainData_trimmed$Class <- factor(trainData_trimmed$Class)

testData_trimmed_1 <- subset(testData_1, select = c(Class,
YearsSinceBuilt, Priceperbuildingarea))
#testData_trimmed$Class <- factor(testData_trimmed$Class)

# Convert the columns to numeric
selected_data <- data_1[, c("Class", "YearsSinceBuilt",
"Priceperbuildingarea")]
selected_data <- sapply(selected_data, as.numeric)

# Check if there are any missing values
if (anyNA(selected_data)) {
```

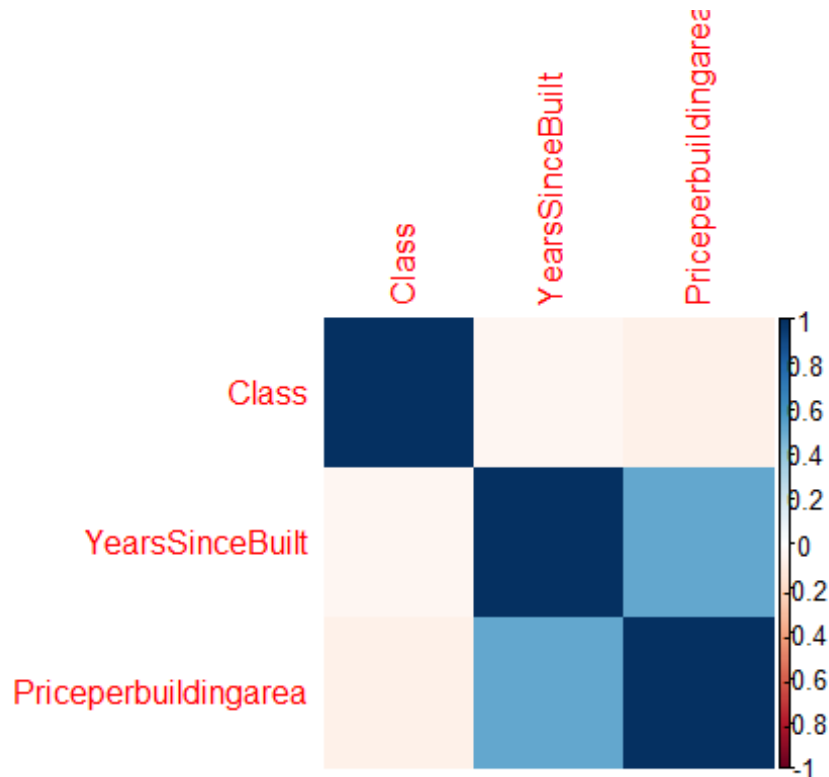
```

# Handle missing values as needed
selected_data <- na.omit(selected_data)
}

# Calculate the correlation matrix
correlation_matrix <- cor(selected_data, use = "pairwise.complete.obs")

# Create the colored correlation grid
corrplot(correlation_matrix, method = "color")

```



### Step 7: Run Regression

```

model_1 <- lm(Priceperbuildingarea ~ ., data = trainData_trimmed_1)
summary(model_1)

##
## Call:
## lm(formula = Priceperbuildingarea ~ ., data = trainData_trimmed_1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9142.2 -1537.9  -169.2  1298.5 10299.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   5354.0584    69.6539   76.867  <2e-16 ***
## Class2         26.4614    72.5400    0.365   0.715

```



```

## Class3          917.8225    79.4400  11.554   <2e-16 ***
## Class4         -2329.5063   114.2920 -20.382   <2e-16 ***
## YearsSinceBuilt  38.3039     0.7665  49.972   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2211 on 6827 degrees of freedom
## Multiple R-squared:  0.3621, Adjusted R-squared:  0.3618
## F-statistic: 968.9 on 4 and 6827 DF,  p-value: < 2.2e-16

# Load necessary Libraries
library(ggplot2)

# Extract residuals and fitted values
residuals_1 <- residuals(model_1)
fitted_values_1 <- fitted(model_1)


# Create a data frame
data_df_1 <- data.frame(residuals_1 = residuals_1, fitted_values_1 =
fitted_values_1)

# QQ Plot
qqplot <- ggplot(data.frame(residuals_1 = residuals_1), aes(sample =
residuals_1)) +
  geom_qq() +
  geom_qq_line() +
  ggtitle("QQ Plot of Residuals") +
  theme_minimal()


# Residual Plot
residual_plot <- ggplot(data_df_1, aes(x = fitted_values_1, y =
residuals_1)) +
  geom_point() +
  geom_hline(yintercept = 0, linetype = "dashed", color = "red") +
  ggtitle("Residual Plot") +
  xlab("Fitted Values") +
  ylab("Residuals") +
  theme_minimal()

# Show plots
print(qqplot)

```



```
print(residual_plot)
```



## Step 8: Show Algorithm Metrics

```
predictions_1 <- predict(model_1, testData_trimmed_1)
actual_1 <- testData_trimmed_1$Priceperbuildingarea

# Calculate RMSE and MAE
RMSE_1 <- sqrt(mean((predictions_1 - actual_1) ^ 2))
MAE_1 <- mean(abs(predictions_1 - actual_1))

# Print metrics
print(paste("RMSE:", RMSE_1))

## [1] "RMSE: 2305.1990809635"

print(paste("MAE:", MAE_1))

## [1] "MAE: 1738.2692701769"
```

## refine the model by introducing distance from center

### Step 9: calculate geo center

```
data_2=data_1

# Calculate the center point
center_longitude <- mean(data_2$Longitude)
center_latitude <- mean(data_2$Latitude)
# Print the center point
cat("Center Longitude:", center_longitude, "\n")

## Center Longitude: 144.9909

cat("Center Latitude:", center_latitude, "\n")

## Center Latitude: -37.80384
```

### Step 10: calculate distance

```
# Function to calculate distance between two points given their
Longitude and Latitude
haversine_distance <- function(lon1, lat1, lon2, lat2) {
  # Convert Latitude and Longitude from degrees to radians
  lon1 <- lon1 * pi / 180
  lat1 <- lat1 * pi / 180
  lon2 <- lon2 * pi / 180
  lat2 <- lat2 * pi / 180

  # Haversine formula
  dlon <- lon2 - lon1
  dlat <- lat2 - lat1
  a <- sin(dlat/2)^2 + cos(lat1) * cos(lat2) * sin(dlon/2)^2
  c <- 2 * asin(sqrt(a))
}
```

```

# Radius of the Earth in kilometers
R <- 6371

# Calculate the distance
distance <- R * c
return(distance)
}

# Calculate the center point
center_longitude <- mean(data_2$Longitude)
center_latitude <- mean(data_2$Latitude)
# Print the center point
cat("Center Longitude:", center_longitude, "\n")

## Center Longitude: 144.9909

cat("Center Latitude:", center_latitude, "\n")

## Center Latitude: -37.80384

# Calculate distance from center for each data point
data_2$distance_from_center <- apply(data_2, 1, function(row) {
  haversine_distance(as.numeric(row["Longitude"]),
as.numeric(row["Latitude"]), center_longitude, center_latitude)
})

# Print the updated data frame
print(data_2)

## # A tibble: 8,537 × 26
##   Suburb Address Rooms Type   Price Method SellerG Date
##   <chr> <chr> <dbl> <chr> <dbl> <chr> <chr> <date>
##   <dbl> <dbl>
## 1 Abbot... 25 Blo...    2 h    1.03e6 S    Biggin 2016-02-04
## 2.5 3067
## 2 Abbot... 5 Char...    3 h    1.46e6 SP   Biggin 2017-03-04
## 2.5 3067
## 3 Abbot... 55a Pa...    4 h    1.6 e6 VB   Nelson 2016-06-04
## 2.5 3067
## 4 Abbot... 124 Ya...    3 h    1.88e6 S    Nelson 2016-05-07
## 2.5 3067
## 5 Abbot... 98 Cha...    2 h    1.64e6 S    Nelson 2016-10-08
## 2.5 3067
## 6 Abbot... 10 Val...    2 h    1.10e6 S    Biggin 2016-10-08
## 2.5 3067
## 7 Abbot... 40 Nic...    3 h    1.35e6 VB   Nelson 2016-11-12
## 2.5 3067
## 8 Abbot... 123/56...    2 u    7.5 e5 S    Biggin 2016-11-12
## 2.5 3067
## 9 Abbot... 16 Wil...    2 h    1.31e6 S    Jellis 2016-10-15

```

```

2.5      3067
## 10 Abbot... 42 Hen...      3 h      1.20e6 S      Jellis  2016-07-16
2.5      3067
## # i 8,527 more rows
## # i 16 more variables: Bedroom2 <dbl>, Bathroom <dbl>, Car <dbl>,
## #   Landsize <dbl>, BuildingArea <dbl>, YearBuilt <dbl>, CouncilArea
## #   <chr>,
## #   Latitude <dbl>, Longitude <dbl>, Regionname <chr>, Propertycount
## #   <dbl>,
## #   YearOfSale <dbl>, YearsSinceBuilt <dbl>, Priceperbuildingarea
## #   <dbl>,
## #   Class <fct>, distance_from_center <dbl>

```

## Step 11: Split Data into Training and Testing

```

library(caret)
set.seed(123) # For reproducibility
index <- createDataPartition(data_2$Priceperbuildingarea, p=0.8,
list=FALSE)
trainData_2 <- data_2[index, ]
testData_2 <- data_2[-index, ]

trainData_trimmed_2=subset(trainData_2, select =
c(distance_from_center, YearsSinceBuilt, Priceperbuildingarea))
#trainData_trimmed$Class <- factor(trainData_trimmed$Class)

testData_trimmed_2=subset(testData_2, select = c(distance_from_center,
YearsSinceBuilt, Priceperbuildingarea))
#testData_trimmed$Class <- factor(testData_trimmed$Class)

```

## Step 12: Run Regression

```

model_2 <- lm(Priceperbuildingarea ~ ., data = trainData_trimmed_2)
summary(model_2)

##
## Call:
## lm(formula = Priceperbuildingarea ~ ., data = trainData_trimmed_2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9173.5 -1389.2  -185.5  1214.2 11771.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    7594.6911    67.6008   112.35 <2e-16 ***
## distance_from_center -146.4118     3.3025   -44.33 <2e-16 ***
## YearsSinceBuilt     31.2631     0.7414    42.17 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2070 on 6829 degrees of freedom

```

```
## Multiple R-squared:  0.441, Adjusted R-squared:  0.4408
## F-statistic:  2693 on 2 and 6829 DF,  p-value: < 2.2e-16

# Extract residuals and fitted values for model2
residuals_2 <- residuals(model_2)
fitted_values_2 <- fitted(model_2)


# Create a data frame for model2
data_df_model_2 <- data.frame(residuals_2 = residuals_2,
fitted_values_2 = fitted_values_2)

# QQ Plot for model2
qqplot_model <- ggplot(data.frame(residuals_2 = residuals_2),
aes(sample = residuals_2)) +
  geom_qq() +
  geom_qq_line() +
  ggtitle("QQ Plot of Residuals - Model 2") +
  theme_minimal()

# Residual Plot for model2
residual_plot_model <- ggplot(data_df_model_2, aes(x = fitted_values_2,
y = residuals_2)) +
  geom_point() +
  geom_hline(yintercept = 0, linetype = "dashed", color = "red") +
  ggtitle("Residual Plot - Model 2") +
  xlab("Fitted Values") +
  ylab("Residuals") +
  theme_minimal()


# Show plots for model1
print(qqplot_model)
```

QQ Plot of Residuals - Model 2



```
print(residual_plot_model)
```


Residual Plot - Model 2



### Step 13: visual comparison of both models

```
# Predicted vs Actual values for model 1(distance based model)  
plot(trainData_trimmed_1$Priceperbuildingarea, predict(model_1),  
      xlab = "Actual Priceperbuildingarea", ylab = "Predicted  
Priceperbuildingarea",  
      main = "Model 1: Predicted vs Actual")  
  
# Add a reference line with slope = 1  
abline(0, 1, col = "red")
```

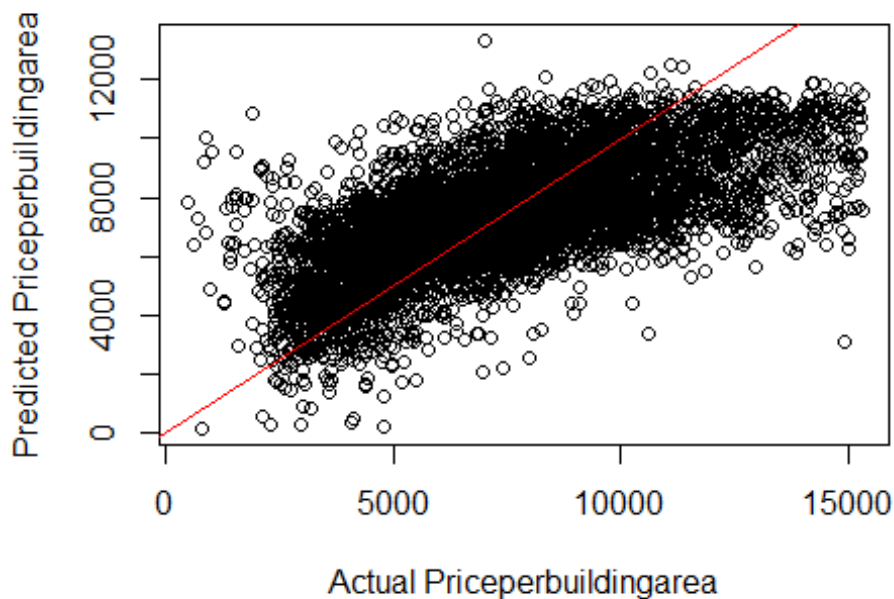
#### Model 1: Predicted vs Actual



```
# Predicted vs Actual values for model 2(class based model)  
plot(trainData_trimmed_2$Priceperbuildingarea, predict(model_2),  
      xlab = "Actual Priceperbuildingarea", ylab = "Predicted  
Priceperbuildingarea",  
      main = "Model 2: Predicted vs Actual")  
  
# Add a reference line with slope = 1  
abline(0, 1, col = "red")
```



## Model 2: Predicted vs Actual




```
# Set up the plotting area
par(mfrow = c(1, 2))


# Predicted vs Actual values for model 1(distance based model)
plot(trainData_trimmed_1$Priceperbuildingarea, predict(model_1),
      xlab = "Actual Priceperbuildingarea", ylab = "Predicted
Priceperbuildingarea",
      main = "Model 1: Predicted vs Actual")
abline(0, 1, col = "red") # Add a reference line with slope = 1

# Predicted vs Actual values for model 2(class based model)
plot(trainData_trimmed_2$Priceperbuildingarea, predict(model_2),
      xlab = "Actual Priceperbuildingarea", ylab = "Predicted
Priceperbuildingarea",
      main = "Model 2: Predicted vs Actual")
abline(0, 1, col = "red") # Add a reference line with slope = 1
```

## Model 1: Predicted vs Actual    Model 2: Predicted vs Actual



Actual Price per building area




Actual Price per building area

```
# Combine data frames for both models
combined_data <- rbind(
  data.frame(Model = "Model_1", residuals = residuals_1, fitted_values
= fitted_values_1),
  data.frame(Model = "Model_2", residuals = residuals_2, fitted_values
= fitted_values_2)
)


# QQ Plot
qqplot_combined <- ggplot(combined_data, aes(sample = residuals, color
= Model)) +
  geom_qq() +
  geom_qq_line() +
  ggtitle("QQ Plot of Residuals") +
  theme_minimal()

# Residual Plot
residual_plot_combined <- ggplot(combined_data, aes(x = fitted_values,
y = residuals, color = Model)) +
  geom_point() +
  geom_hline(yintercept = 0, linetype = "dashed", color = "red") +
  ggtitle("Residual Plot") +
  xlab("Fitted Values") +
  ylab("Residuals") +
  theme_minimal()
```

```
# Show combined plots  
print(qqplot_combined)
```




```
print(residual_plot_combined)
```



```
# Calculate lower and upper bounds for residuals (e.g., 95% confidence interval)
lower_bound <- quantile(combined_data$residuals, 0.025)
upper_bound <- quantile(combined_data$residuals, 0.975)

# Residual Plot with bounds
residual_plot_combined <- ggplot(combined_data, aes(x = fitted_values,
y = residuals, color = Model)) +
  geom_point() +
  geom_hline(yintercept = 0, linetype = "dashed", color = "red") +
  geom_hline(yintercept = lower_bound, linetype = "dashed", color =
"blue") +
  geom_hline(yintercept = upper_bound, linetype = "dashed", color =
"blue") +
  ggtitle("Residual Plot with Bounds") +
  xlab("Fitted Values") +
  ylab("Residuals") +
  theme_minimal()

# Show the residual plot with bounds
print(residual_plot_combined)
```



```

library(gridExtra)

##
## 载入程辑包: 'gridExtra'

## The following object is masked from 'package:dplyr':
##
##   combine

# Calculate lower and upper bounds for residuals (e.g., 95% confidence
# interval)
lower_bound <- quantile(combined_data$residuals, 0.025)
upper_bound <- quantile(combined_data$residuals, 0.975)

# QQ Plot for Model 1
qqplot_model1 <- ggplot(combined_data[combined_data$Model == "Model_1",
], aes(sample = residuals)) +
  geom_qq() +
  geom_qq_line() +
  ggtitle("Model 1: QQ Plot of Residuals") +
  theme_minimal()

# Residual Plot for Model 1
residual_plot_model1 <- ggplot(combined_data[combined_data$Model ==
"Model_1", ], aes(x = fitted_values, y = residuals)) +
  geom_point() +
  geom_hline(yintercept = 0, linetype = "dashed", color = "red") +

```

```

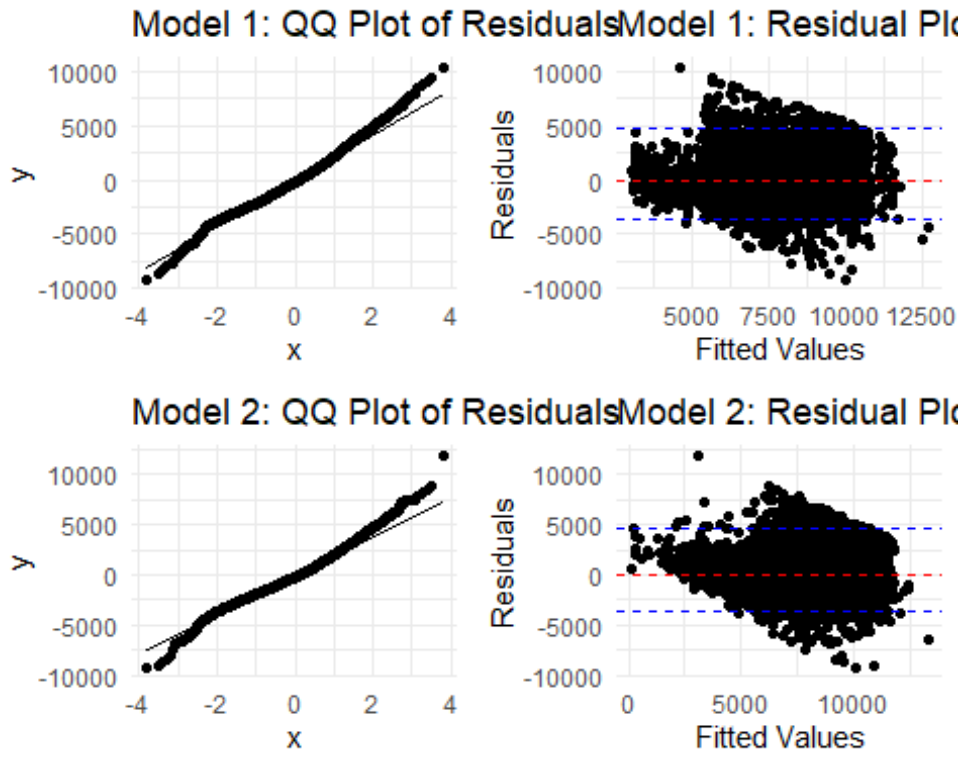
  geom_hline(yintercept = lower_bound, linetype = "dashed", color =
"blue") +
  geom_hline(yintercept = upper_bound, linetype = "dashed", color =
"blue") +
  ggtitle("Model 1: Residual Plot") +
  xlab("Fitted Values") +
  ylab("Residuals") +
  theme_minimal()

# QQ Plot for Model 2
qqplot_model2 <- ggplot(combined_data[combined_data$Model == "Model_2",
], aes(sample = residuals)) +
  geom_qq() +
  geom_qq_line() +
  ggtitle("Model 2: QQ Plot of Residuals") +
  theme_minimal()

# Residual Plot for Model 2
residual_plot_model2 <- ggplot(combined_data[combined_data$Model ==
"Model_2", ], aes(x = fitted_values, y = residuals)) +
  geom_point() +
  geom_hline(yintercept = 0, linetype = "dashed", color = "red") +
  geom_hline(yintercept = lower_bound, linetype = "dashed", color =
"blue") +
  geom_hline(yintercept = upper_bound, linetype = "dashed", color =
"blue") +
  ggtitle("Model 2: Residual Plot") +
  xlab("Fitted Values") +
  ylab("Residuals") +
  theme_minimal()

# Arrange plots in a 2x2 grid
grid.arrange(qqplot_model1, residual_plot_model1, qqplot_model2,
residual_plot_model2, ncol = 2, nrow = 2)

```



### Step 14: Comparison of metrics

```
summary(model_1)

##
## Call:
## lm(formula = Priceperbuildingarea ~ ., data = trainData_trimmed_1)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9142.2 -1537.9  -169.2  1298.5 10299.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  5354.0584    69.6539   76.867 <2e-16 ***
## Class2       26.4614    72.5400    0.365  0.715
## Class3       917.8225    79.4400   11.554 <2e-16 ***
## Class4     -2329.5063   114.2920  -20.382 <2e-16 ***
## YearsSinceBuilt  38.3039    0.7665   49.972 <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2211 on 6827 degrees of freedom
## Multiple R-squared:  0.3621, Adjusted R-squared:  0.3618
## F-statistic: 968.9 on 4 and 6827 DF,  p-value: < 2.2e-16

summary(model_2)
```

```

##
## Call:
## lm(formula = Priceperbuildingarea ~ ., data = trainData_trimmed_2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -9173.5 -1389.2  -185.5  1214.2 11771.7
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)    7594.6911     67.6008  112.35  <2e-16 ***
## distance_from_center -146.4118      3.3025  -44.33  <2e-16 ***
## YearsSinceBuilt     31.2631      0.7414   42.17  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 2070 on 6829 degrees of freedom
## Multiple R-squared:  0.441, Adjusted R-squared:  0.4408
## F-statistic: 2693 on 2 and 6829 DF, p-value: < 2.2e-16

# Metrics for model 1
metrics_model_1 <- summary(model_1)
r_squared_model_1 <- metrics_model_1$r.squared
adj_r_squared_model_1 <- metrics_model_1$adj.r.squared
residual_std_error_model_1 <- sqrt(metrics_model_1$sigma^2)
f_statistic_model_1 <- metrics_model_1$fstatistic[1]

# Metrics for model 2
metrics_model_2 <- summary(model_2)
r_squared_model_2 <- metrics_model_2$r.squared
adj_r_squared_model_2 <- metrics_model_2$adj.r.squared
residual_std_error_model_2 <- sqrt(metrics_model_2$sigma^2)
f_statistic_model_2 <- metrics_model_2$fstatistic[1]

# Print metrics for both models
cat("Model 1 Metrics:\n")

## Model 1 Metrics:

cat("R-squared:", r_squared_model_1, "\n")

## R-squared: 0.3621294

cat("Adjusted R-squared:", adj_r_squared_model_1, "\n")

## Adjusted R-squared: 0.3617557

cat("Residual Standard Error:", residual_std_error_model_1, "\n")

## Residual Standard Error: 2211.019

cat("F-statistic:", f_statistic_model_1, "\n\n")

```



```
## F-statistic: 968.9494
cat("Model 2 Metrics:\n")
## Model 2 Metrics:
cat("R-squared:", r_squared_model_2, "\n")
## R-squared: 0.4409704
cat("Adjusted R-squared:", adj_r_squared_model_2, "\n")
## Adjusted R-squared: 0.4408067
cat("Residual Standard Error:", residual_std_error_model_2, "\n")
## Residual Standard Error: 2069.569
cat("F-statistic:", f_statistic_model_2, "\n")
## F-statistic: 2693.406
```

**Step 15: Final comment**

**Step 16: Business application**